

RBE595 - Week 2 Programming Assignment

Keith Chester

Due date: January 14, 2023

This is a programming assignment where we were tasked with recreating a figure in our textbooks. The figure demonstrated the effects of modifying the ϵ for epsilon-greedy approaches. That figure and the code utilized to generate it is shared below:




Figure 1: Recreated figure

```
from math import sqrt
from random import gauss, randint, uniform
from typing import List, Tuple

import matplotlib.pyplot as plt

def perform_run(
    steps: int, epsilon: float, k_bandits: int
) -> Tuple[List[float], List[int]]:
    """
    perform_run takes a given number of steps and an epsilon, and then takes
    that run for :steps: steps. Once this is done it returns a list of each
    rewards on that step, and whether or not it was the optimal action
    (represented as 0 or 1)
    """
    rewards: List[float] = []
```

```

optimal: List[int] = []

# We need to create a specified amount of bandits with mean 0
# and variance 1 for our problem
mean = 0
variance = 1
bandits = [gauss(mean, sqrt(variance)) for i in range(0, k_bandits)]
best_bandit = bandits.index(max(bandits))
# We assume some knowledge about the bandits. Over time we will learn
# more as we explore.
knowledge = [gauss(mean, sqrt(variance)) for i in range(0, k_bandits)]

for i in range(0, steps):
    if uniform(0, 1) <= 1 - epsilon:
        # In this situation we act greedily.
        index = knowledge.index(max(knowledge))
    else:
        # We explore here
        index = randint(0, k_bandits - 1)

    # Grab our reward
    reward = bandits[index]
    knowledge[index] = reward
    rewards.append(reward)
    # Is it optimal? Record that
    optimal.append(1 if index == best_bandit else 0)

return rewards, optimal

def run_experiment(runs: int, steps_per_run: int, epsilon: float, k_bandits: int):
    all_rewards: List[float] = [0.0] * steps_per_run
    all_optimal: List[float] = [0.0] * steps_per_run

    for _ in range(0, runs):
        rewards, optimal = perform_run(steps_per_run, epsilon, k_bandits)

        # We are iteratively calculating the all_rewards
        # and all_optimal averages so we can avoid holding
        # it all in memory
        for index, _ in enumerate(all_rewards):
            all_rewards[index] += rewards[index] / runs
            all_optimal[index] += optimal[index] / runs

    return all_rewards, all_optimal

if __name__ == "__main__":
    runs = 2000
    steps = 1000
    bandits = 10
    rewards_00, optimal_00 = run_experiment(runs, steps, 0.00, bandits)
    rewards_01, optimal_01 = run_experiment(runs, steps, 0.01, bandits)
    rewards_10, optimal_10 = run_experiment(runs, steps, 0.10, bandits)

    figure, axis = plt.subplots(2, sharex=True)
    plt.xlabel("Steps")
    figure.suptitle("Greedy k-Bandits")

    axis[0].set_title("Average reward")
    axis[0].plot(rewards_00, label=" =0")
    axis[0].plot(rewards_01, label=" =.01")
    axis[0].plot(rewards_10, label=" =.1")
    axis[0].legend(loc="lower center")

    axis[1].set_title("% Optimal Action")
    axis[1].plot(optimal_00, label=" =0")
    axis[1].plot(optimal_01, label=" =0.01")
    axis[1].plot(optimal_10, label=" =0.1")
    axis[1].legend(loc="lower center")

```

```
figure.savefig("./epsilon_greedy.png")
```