RBE595 - Week 10 Assignment

Keith Chester

Due date: March 19, 2023

Problem 1

What is "planning" in the context of Reinforcement Learning?

"Planning" in this context is the process of action selection in an environment while aiming to maximize a reward. Alternatively, "learning" involves updating a value and/or policy function through observed experience.

Problem 2

What is the difference between Dyna-Q and Dyna-Q+ algorithms?

There are two notable differences. The first is that Dyna-Q+ includes an exploration bonus term in its update rule for its Q function. This results in an agent that is encouraged to explore states/actions that haven't been explored yet, which is beneficial in non-static enviornments especially.

Another notable difference is that Dyna-Q+ utilizes a count-based mechanism to estimate the "novelty" of a state/action pair; this is used to calculate the aforementioned exploration bonus.

Problem 3

Model-based RL methods suffer more bias than model-free methods. Is this statement correct? Why or why not?

This statement isn't correct. Model-free RL methods learn on-line actively interacting with the environment (which is necessary given the scenario, such as an environment that is too complex to simulate or is not understood at all). Alternatively model-based RL methods learn a model of an environment (or attempt to build a representation as best as possible) and from this tries to simulate the outcomes of different actions prior to actually interacting with the environment. If the model itself is a poor representation of the environment then it will introduce bias. If the model is a perfect or good-enough representation of the environment, however, this can be more efficient and less biased than an equivalent model-free approach.

Problem 4

Model-based RL methods are more sample efficient. Is this statement correct? Why or why not?

Yes, this statement is generally correct, depending on your environment and the accuracy of your representation of the environment. This is because you can quickly simulate additional training data/experience throug the simulated environment and thus lowering the cost to sample the environment, as opposed to an agent that has to explore a real-world environment.

Problem 5

What are the 4 steps of the MCTS algorithm? How does MCTS balances exploration/exploitation?

The four steps of the Monte Carlo Tree Search (MCTS) algorithm are selection, expansion, simulation, and backup.

Selection - Starting from the root node of the tree, a path is traversed through the tree to a leaf node using a selection polciy. The selection policy balances exploration/exploitation at this stage by choosing either nodes with high value or high uncertainty (exploitation and exploration, respectively).

Expanaion - On some iterations we create new children nodes added to the leaf node when the selection policy chooses unexplored options.

Simulation - Starting from the selected or resulting new child node we simulate a complete episode with the rollout policy.

Backup - The return generated by this episode is propagated ("backed up") through the tree updating its calculated value estimations.

Problem 6

The nonplanning method looks particularly poor in Figure 8.3 because it is a one-step method; a method using multi-step bootstrapping would do better. Do you think one of the multi-step bootstrapping methods from Chapter 7 could do as well as the Dyna method? Explain why or why not

A multi-step bootstrapping method (MSBM) likely will not do as well as the Dyna method. MSBM have to experience the environment to learn its value estimation function, but it's model-free and thus can't leveraged the learned model of the envrionment in order to bootstrap with additional training data. The Dyna method, by ocntrast, has higher sample-efficiency and is generally more effective in complex and/or stochastic and/or changing environments.

Problem 7

Why did the Dyna agent with exploration bonus, Dyna-Q+, perform better in the first phase as well as in the second phase of the blocking and shortcut experiments?

The exploration bonus encourages the agent to visit states that were previously unvisited. This increased emphasis on exploration means the agent has a better idea of the environment and can better find optimal solutions to a given environment. Dyna in contrast doesn't explore as much and can become stuck on a set solution in examples where the environment is stochastic.